

Deriving MLE estimates for a general multinomial distribution

1/1 point (graded)

In this problem, we will derive the maximum likelihood estimates for a multinomial distribution with more than 2 parameters. For this we employ a powerful optimization strategy called method of lagrange multipliers.

First let $P(D|\theta)$ denote the probability of a multinomial model M with parameters $\theta = \{\theta_1, \theta_2 \dots \theta_N\}$ generating a document D .

Which of the following option lists the correct expression for $P(D|\theta)$. Choose from options below.

☐ $P(D|\theta) = \sum_{w \in W} \theta_w^{\text{count}(w)}$

☒ $P(D|\theta) = \prod_{w \in W} \theta_w^{\text{count}(w)}$ ✓

☐ $P(D|\theta) = \prod_{w \in W} \text{count}(w) \theta_w$

☐ $P(D|\theta) = \prod_{w \in W} \theta_w + \text{count}(w)$

Solution:

Recall from the lecture that

$$P(D|\theta) = \prod_{i=1}^n \theta_{w_i} = \prod_{w \in W} \theta_w^{\text{count}(w)}$$

Submit

You have used 1 of 2 attempts

i Answers are displayed within the problem

Constraints

1/1 point (graded)

Which of the following options lists the right set of constraints on the parameters θ_w of the multinomial model.

☒ $\theta_w \geq 0, \sum_{w \in W} \theta_w = 1$ ✓

☐ $\theta_w \geq 0, \sum_{w \in W} \theta_w < 1$

☐ $\theta_w < 0, \sum_{w \in W} \theta_w > -1$

☐ $\theta_w \geq 0, \sum_{w \in W} \theta_w \geq 1$

Solution:

Note that θ_w denotes the probability of model M choosing the word w . Since it's a probability, its value must lie between 0 and 1. Therefore, $0 \leq \theta_w \leq 1$.

Further, all the above probability values must also sum up to 1. That is, $\sum_{w \in W} \theta_w = 1$.

Submit

You have used 1 of 2 attempts

i Answers are displayed within the problem

Stationary points for lagrange function

2/2 points (graded)

Let's recall the function that we're trying to optimize:

$$\max_{\theta} P(D|\theta) = \max_{\theta} \prod_{w \in W} \theta_w^{\text{count}(w)}$$

Let us take natural logarithm on both sides of the equation in order to bring down the exponent

$$\max_{\theta} \log P(D|\theta) = \max_{\theta} \sum_{w \in W} \text{count}(w) \log \theta_w$$

subject to the following constraints,

$$\sum_{w \in W} \theta_w = 1$$

Let's define a function L called the lagrange function for the sake of the above defined constrained optimization problem:

$$L = \log P(D|\theta) + \lambda (\sum_{w \in W} \theta_w - 1)$$

where λ is a constant.

Consider finding stationary points for L by solving for equation obtained by setting its derivative to zero. That is,

$$\frac{\partial}{\partial \theta_w}(\log P(D|\theta) + \lambda (\sum_{w \in W} \theta_w - 1)) = 0$$

Solve for θ_w from the above equation. Choose the right answer for θ_w from options below.

- ☐
 $\theta_w = \frac{-\lambda}{count(w)}$
- ☐
 $\theta_w = \lambda \times count(w)$
- ☐
 $\theta_w = -\lambda \times count(w)$
- ☒
 $\theta_w = \frac{-count(w)}{\lambda}$ ✓

Now, apply the constraint that $\sum_{w \in W} \theta_w = 1$, we get that λ is:

- ☒
 $\lambda = -\sum_{w \in W} count(w)$ ✓
- ☐
 $\lambda = \sum_{w \in W} count(w)$
- ☐
 $\lambda = -\sum_{w \in W} count(w) \times \theta_w$
- ☐
 $\lambda = \sum_{w \in W} count(w) \times \theta_w$

Solution:

$$\frac{\partial}{\partial \theta_w}(\log P(D|\theta) + \lambda (\sum_{w \in W} \theta_w - 1)) = 0$$

$$\frac{\partial \log P(D|\theta_w)}{\partial \theta_w} + \lambda = 0$$

$$\frac{\partial \log \Pi_{w \in W} \theta_w^{count(w)}}{\partial \theta_w} + \lambda = 0$$

$$\frac{\partial \sum_{w \in W} \log \theta_w \times count(w)}{\partial \theta_w} + \lambda = 0$$

$$\frac{count(w)}{\theta_w} + \lambda = 0$$

$$\theta_w = -\frac{\text{count}(w)}{\lambda}$$

If we apply the constraint that $\sum_{w \in W} \theta_w = 1$ we get

$$\sum_{w \in W} \theta_w = 1$$

$$\sum_{w \in W} -\frac{\text{count}(w)}{\lambda} = 1$$

$$\sum_{w \in W} \text{count}(w) = -\lambda$$

$$\lambda = -\sum_{w \in W} \text{count}(w)$$

Substituting this expression for λ back into our previous expression for θ_w we get

$$\theta_w = -\frac{\text{count}(w)}{\lambda}$$

$$\theta_w = \frac{\text{count}(w)}{\sum_{w \in W} \text{count}(w)}$$

Note that $\theta_w > 0$ and $\sum_{w \in W} \theta_w = 1$ satisfying all the constraints. These set of θ_w parameters are the maximum likelihood estimates for this multinomial generative distribution.

Submit

You have used 1 of 2 attempts

i Answers are displayed within the problem

Discussion

Show Discussion

Topic: Unit 4 Unsupervised Learning (2 weeks) :Lecture 15. Generative Models / 6. MLEs for general multinomial distribution