2. Learning to Control: Introduction to Reinforcement Learning

Learning to Control: Introduction Reinforcement Learning (RL)

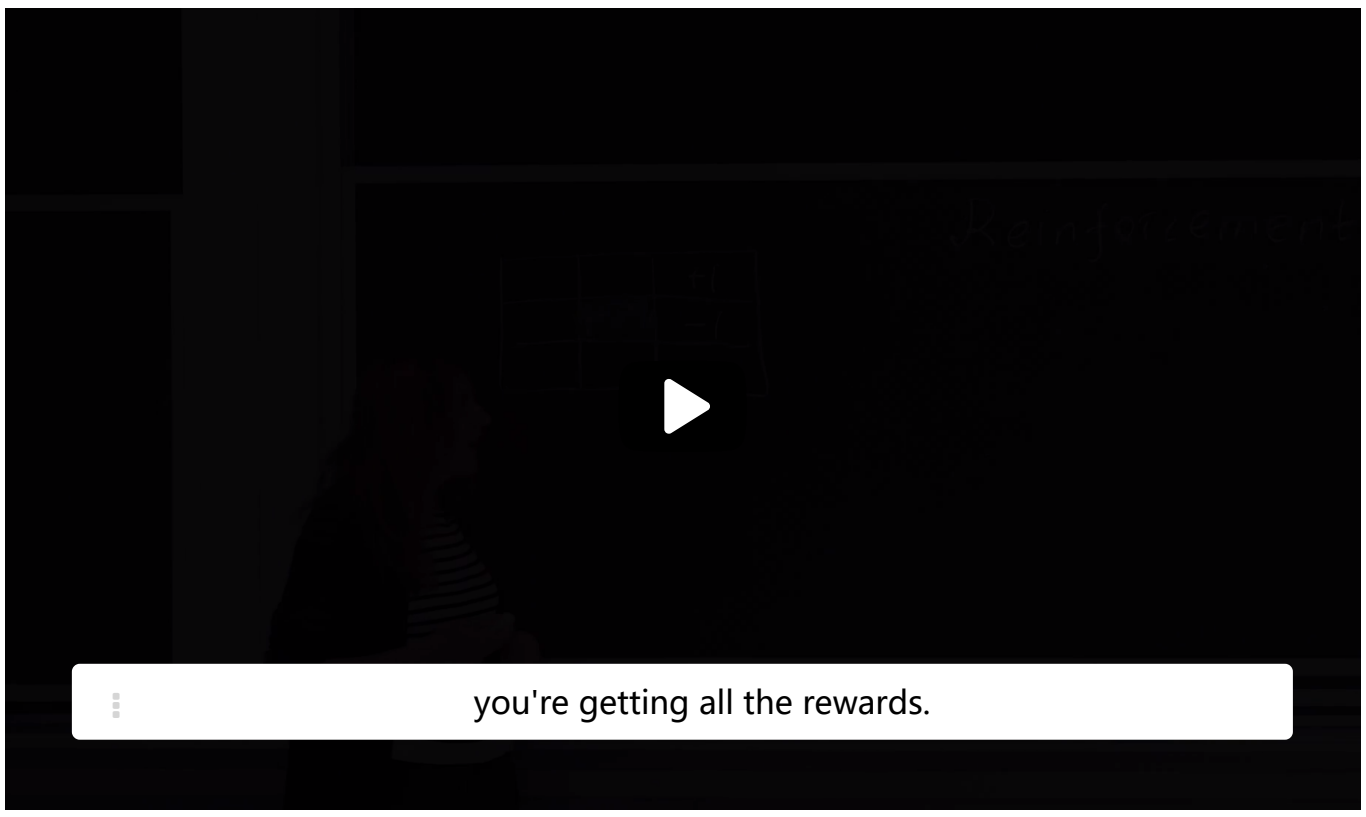

And if the robot did the good job,

it would eventually end up here and get reward plus 1.

And if the robot fail--

not fail, but if the robot end up here,

it's going to get a negative reward or punishment, OK?

So your goal is to kind of find to find a strategy

to get to this particular place where

you're getting all the rewards.

End of transcript. Skip to the start.

Video
Download video file

Transcripts
Download SubRip (.srt) file
Download Text (.txt) file

Objectives of RL

1/1 point (graded)
Select one or more of the statements that are true about Reinforcement Learning (RL) from the options below:

- [ ] The goal of RL is to minimize the loss between predictions and labels for points in a labelled training dataset
- [x] The goal of RL is to learn a good policy with none or limited supervision ✔
- [ ] For an RL algorithm to work, it must a receive a non-zero reward after every step
- [ ] RL works by maximizing the reward for each immediate next step

✔

Solution:

- Unlike with Supervised Learning, there would typically be no labelled training dataset associated with Reinforcement Learning tasks. RL algorithms learn to pick “good" actions based on the rewards that they receive during training.
- RL is most applicable to tasks where there is no clear cut supervised training data available.

Reinforcement learning algorithms can learn to take actions so as to maximize some notion of a cumulative reward instead of the reward for the next step and they can take “good" actions even without any intermediate rewards.