

Reinforcement Learning: Time for Action!

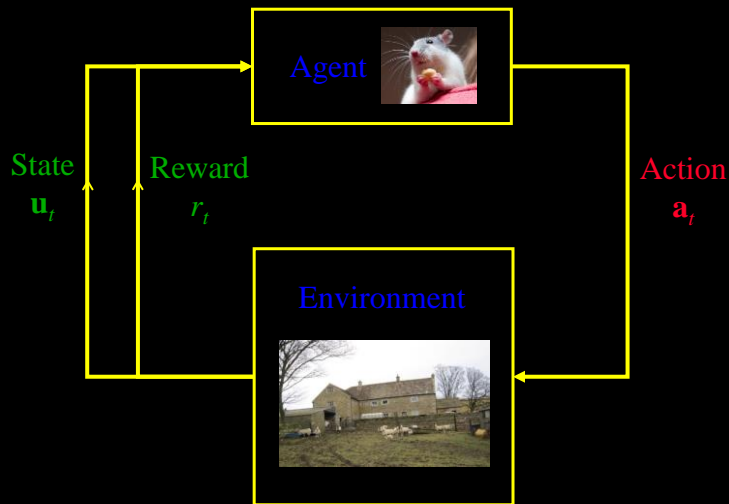



Image Source: Wikimedia Commons

The Problem




Learn a state-to-action
mapping or “policy”:

$$\pi(\mathbf{u}) = \mathbf{a}$$

which maximizes the
expected total future
reward:

$$\left\langle \sum_{\tau=0}^{T-t} r(t+\tau) \right\rangle_{\text{trials}}$$

Example: Rat in a barn




States = locations A, B, or C

Actions= L (go left) or R (go right)

If the rat chooses L or R *at random* (random “policy”), what is the expected reward (or “value”) v for each state?

Image Source: Dayan & Abbott textbook

Policy Evaluation



For random policy:

$$v(B) = \frac{1}{2} \cdot 0 + \frac{1}{2} \cdot 5 = 2.5$$

$$v(C) = \frac{1}{2} \cdot 2 + \frac{1}{2} \cdot 0 = 1$$

$$v(A) = \frac{1}{2} \cdot v(B) + \frac{1}{2} \cdot v(C) = 1.75$$

ONLINE?

Let value of state u
 $v(u)$ = weight $w(u)$

Can learn value of states using TD learning:

$$w(u) \leftarrow w(u) + \varepsilon [r(u) + v(u') - v(u)]$$


(Location, action) \Rightarrow new location i.e., $(u, a) \Rightarrow u'$

现在的reward

下一步后的期望reward

这一步的期望reward

TD Learning of Values for Random Policy



(For all three,
 $\epsilon = 0.5$)


Once I know the values, I can pick the action
that leads to the higher valued state!



5

Image Source: Dayan & Abbott textbook

Selecting Actions based on Values



Values act as
surrogate immediate
rewards → Locally
optimal choice leads
to globally optimal
policy for “Markov”
environments
(Related to Dynamic
Programming)

6

Putting it all together: Actor-Critic Learning

- Two separate components: Actor (selects action and maintains policy) and Critic (maintains value of each state)

1. Critic Learning ("Policy Evaluation"):

Value of state $u = v(u) = w(u)$

$$w(u) \leftarrow w(u) + \varepsilon [r(u) + v(u') - v(u)] \quad (\text{same as TD rule})$$

2. Actor Learning ("Policy Improvement"):

$$P(a; u) = \frac{\exp(\beta Q_a(u))}{\sum_b \exp(\beta Q_b(u))}$$

\leftarrow SOFTMAX \leftarrow EXPLORE
 Probabilistically select an action a at state u


For all actions a' :

$$Q_{a'}(u) \leftarrow Q_{a'}(u) + \varepsilon [r(u) + v(u') - v(u)] (\delta_{aa'} - P(a'; u))$$


3. Repeat 1 and 2

7

Actor-Critic Learning in our Barn Example



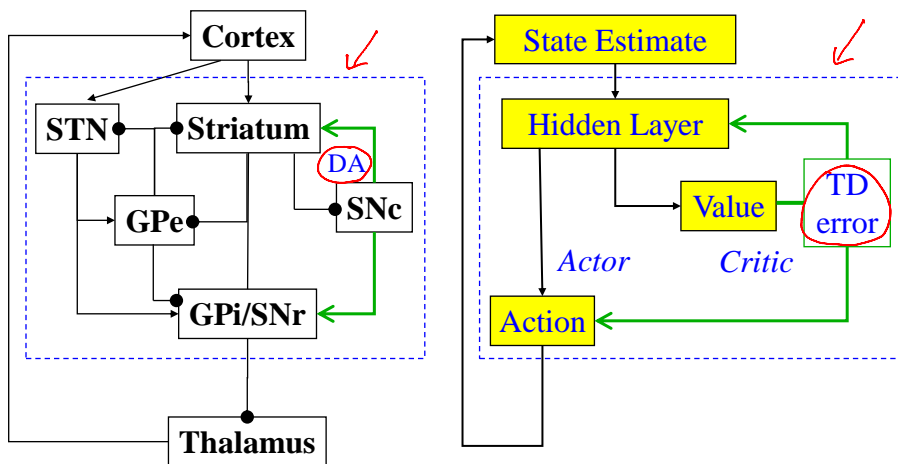
Probability of going Left at each location



8

Image Source: Dayan & Abbott textbook

Possible Implementation of the Actor-Critic Model in the Basal Ganglia



9
(See Supplementary Materials for references)

Reinforcement learning has been applied to
many real-world problems!


Example:

Autonomous Helicopter Flight
(learned from human demonstrations)

↓
(Videos and papers at: <http://heli.stanford.edu/>)

10

Computational Neuroscience



Rajesh P. N. Rao

Adrienne Fairhall

University of Washington, Seattle, USA